

Spatio-temporal modelling and simulation with the FCPP Aggregate Programming framework

I - Introduction

Volker Stolz,[†] Ferruccio Damiani,^{*} Giorgio Audrito^{*}

^{*}University of Turin, Italy

[†]Western Norway University of Applied Sciences, Bergen

June 20, 2022




Outline

- What is Aggregate Computing about?
 - Which computing systems?
 - Which frameworks?
 - What kind of problems?
- Why care about Aggregate Computing?
- So, what Aggregate Computing is?
- How can we formalise it?

Which computing systems?

sequential



parallel

communication by shared memory

usually synchronous

reliable

PRAM model

distributed

communication by messages

usually asynchronous

unreliable

event model



concurrent

Which frameworks?

| | centralised | decentralised |
|----------------------|--|----------------------|
| <i>approach</i> | few special central units responsible of computing and deciding, most units only sensing and actuating | all units are equal |
| <i>collaboration</i> | easy | hard |
| <i>robustness</i> | lower | higher |
| <i>privacy</i> | lower | higher |

AC is a decentralised framework

Which frameworks?

| | “query-based” | “data-based” |
|-----------------------|---|---|
| <i>approach</i> | computing required results whenever queries are issued | maintaining relevant data always updated and locally accessible |
| <i>requirements</i> | none | predictable queries |
| <i>resource usage</i> | on demand (best for fewer queries) | constant (best for frequent queries) |
| <i>responsiveness</i> | lower: requires routing | higher: data locally available |
| <i>quality</i> | higher: fresh results | lower: older results |

AC is designed for “data-based” frameworks
(extensions handling query-based frameworks are possible)

What kind of problems?

(not only) spatial computing problems

position and time of devices are essential inputs of the computation



distance estimation, data summarisation (event detection),
selecting areas (network partitioning, channel establishment...),
inducing shapes (crowd dispersion, formation control...)... and others!

So, what is Aggregate Computing about?

- Which computing systems? → distributed systems
- Which frameworks? → decentralized, (mostly) data-based
- What kind of problems? → born for (but not limited to) spatial computing



is a recent topic. . .

- **cons:** not standard, not well-known, small research group
- **pros:** plenty of things where research is needed

Outline

- What is Aggregate Computing about?
- Why care about Aggregate Computing?
 - Why care for distributed systems at all?
 - Why are distributed systems hard to deal with?
 - How Aggregate Computing can help?
- So, what Aggregate Computing is?
- How can we formalise it?

Why care for distributed systems at all?

- increasing availability of wearable / mobile / embedded / flying devices
- increasing availability of heterogeneous wireless connectivity
- increasing availability of computational resources (device/edge/cloud)
- increasing production/analysis of data, everywhere, anytime



Why are distributed systems hard to deal with?

diverse heterogeneous entities

- different computing power
- sensing and actuation capabilities



we need...

- device **abstraction**
- multi-platform **frameworks**
- not too bad so far...

Why are distributed systems hard to deal with?

data security and privacy

- localised data may be highly personal
- cryptography has got it covered, but hacker attacks do happen
- can we really trust big companies having our data?



in a decentralised system, we can...

- keep our data **locally**
- where communication is needed, share **locally aggregated** neighbour data

Why are distributed systems hard to deal with?

predictable code re-use

- can break down complexity of problems
- needs to avoid unwanted interferences



in a sequential system...

- machine code is hard to re-use
- **abstractions**: objects, mixins...
- all boil down to **functions** and composition

in a distributed system...

- need to **match** messages
- not good with **process calculi**
- ... how to handle multiple instances of same algorithm?

Why are distributed systems hard to deal with?

dynamic goals and environment

- low-end devices may be unreliable...
- as low-power wireless communications between them
- environmental factors may affect the computation



in a centralised system...

- changes in edge devices OK
- changes in central units BAD
- connection loss BAD

in a decentralised system...

- any change OK but HARD
- need re-usable algorithms...
- with precise adaptivity guarantees

Why are distributed systems hard to deal with?

collaboration vs selfishness

- many devices and goals to be achieved
- competitive behaviour may prevent to achieve them
- much easier in centralised systems



in a decentralised system. . .

- we want to achieve a given **global** behaviour,
- by programming **local** interactions
- seek for an automatic **global to local** translation

How Aggregate Computing can help?

challenges are...

- diverse **heterogeneous** entities
- data security and **privacy**
- predictable code **re-use**
- **dynamic** goals and environment
- **collaboration** vs selfishness

aggregate computing provides...

- a **multi-platform** framework
- ... which is **decentralized**
- **composable** abstraction of **behaviour**
- ... and **guarantees**
- with **global-to-local** translation

Outline

- What is Aggregate Computing about?
- Why care about Aggregate Computing?
- So, what Aggregate Computing is?
 - A computation model for proximity-based networks
 - What it provides?
 - Properties
 - How developing works?
- How can we formalise it?

A computation model for proximity-based networks

Network of (possibly mobile) devices

- dynamic topology — induced by (physical or logical) proximity of devices
- each device can receive/send message to/from devices in the neighbourhood

where each device executes a computation/coordination service S

To execute S , each device (asynchronously periodically) **fires**:

- ① collects local (sensor, DB,...) data and received messages
- ② executes a program — **the same for all devices**¹
- ③ sends messages (and performs some actuation)

¹ Supports **integrated development of every part of a distribute application** – cf. *macroprogramming* and *multi-tier programming*:

- R. Newton and M. Welsh (2004). *Region streams: Functional macroprogramming for sensor networks*. In: 1st Workshop on Data Management for Sensor Networks (MSN 2004), pp. 78–87.

<https://doi.org/10.1145/1052199.1052213>

- P. Weisenburger, M. Köhler, G. Salvaneschi (2018). *Distributed system development with Scalaloci*. In: ACM Programming Languages, 2 (OOPSLA):129:1–129:30, 2018. <https://doi.org/10.1145/3276499>

What it provides?

1. Formal language and models for studying algorithms and properties

$$e ::= x \mid \phi \mid c(\bar{e}) \mid b \mid d \mid (\bar{x}) \Rightarrow e \mid e(\bar{e}) \mid \text{nbr}\{e\} \mid \text{old}(e)\{x \Rightarrow e\} \dots$$

2. Language implementations (Proto, Protelis, ScaFi, FCPP)

```
DEF() double abf(ARGS, bool source) { CODE
  return nbr(CALL, INF, [&] (field<double> d) {
    double v = source ? 0.0 : INF;
    return min_hood(CALL, d + node.nbr_dist(), v);
  });
}
```

3. Network simulator (Alchemist, FCPP) or real deployment (FCPP)



Properties

- **Self-stabilisation**: computation resilient to changes in devices state and network topology (behaviour is guaranteed to eventually attain a correct and stable final state despite any transient perturbation in state or topology)¹
- **Eventual consistency**: computation independent from device number and density (behaviour converges to a final state that approximates a predictable limit, based on the continuous environment, as the density and speed of devices increases)²
- **Certified error bounds**: how computation reacts to changes (linking the quality of the services to the amount of computing resources dedicated)³
- ...

¹ M. Viroli, G. Audrito, J. Beal, F. Damiani, D. Pianini (2018). *Engineering resilient collective adaptive systems by self-stabilisation*. ACM Transactions on Modeling and Computer Simulation, 28(2), 2018.

<https://doi.org/10.1145/3177774>

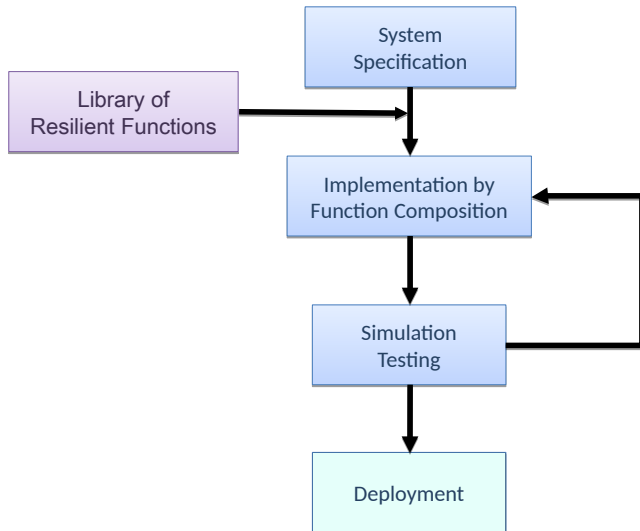
² J. Beal, M. Viroli, D. Pianini, F. Damiani (2017). *Self-adaptation to device distribution in the internet of things*. ACM Transactions on Autonomous and Adaptive Systems 12(3), 2017.

<https://doi.org/10.1145/3105758>

³ G. Audrito, F. Damiani, M. Viroli, E. Bini (2018). *Distributed Real-Time Shortest-Paths Computations with the Field Calculus*. In: IEEE 39th Real-Time Systems Symposium (RTSS 2018).

<https://doi.org/10.1109/RTSS.2018.00013>

How developing works?



Outline


- What is Aggregate Computing about?
- Why care about Aggregate Computing?
- So, what Aggregate Computing is?
- How can we formalise it?
 - What physics can say
 - Event structures
 - Space-time values
 - Space-time functions

What physics can say




Galileo, help!

- events are pairs position, time $\langle \vec{x}, t \rangle$
- position is relative, but time is **universal**
- past, present and future are **linear**




sounds like a sequential/parallel system...

What physics can say



Einstein, help!

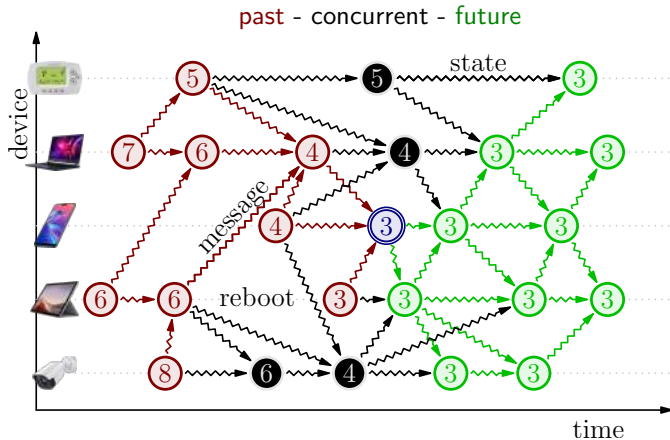
- position and time are both **relative**
- no clock, only **relation** between events
- some events are just **not comparable**



sounds weird just like a distributed system!


Event structures

- a set of **events** E
- a DAG of **messages** \rightsquigarrow
- a **causality** partial order $<$ (transitive closure of \rightsquigarrow)



Space-time values

- maps from events in event structures to values $\Phi : E \rightarrow X$

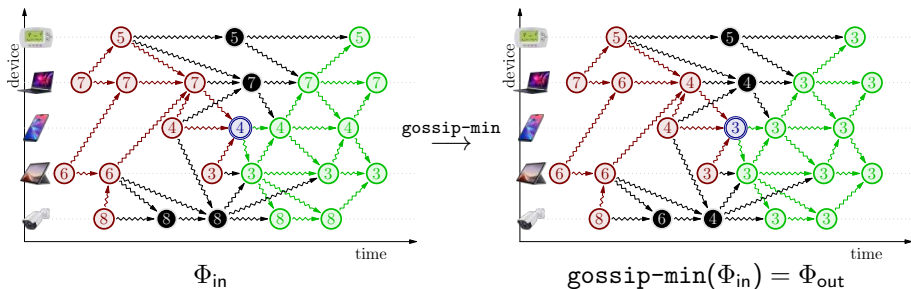


↑ space-time value resulting from a “gossip-min”

every device gossips the minimum of an *input value* and neighbours' *gossips*

Space-time functions

- maps from space-time values to space-time values $f(\Phi_{in}, \dots) = \Phi_{out}$



- they're just plain old functions (maybe with weird domain)
- we can compose them as usual $f(g(\Phi_1), h(\Phi_2), \Phi_3)$